

Intro

The [Floyd-Warshall algorithm](#) for computing all pairwise shortest path lengths in a graph has a computational pattern much like the one for Gaussian elimination. There is a closely related algorithm which is slightly more expensive – $O(n^3 \log n)$ time in general rather than the $O(n^3)$ time required by Floyd-Warshall – but which looks very much like matrix multiplication. In this assignment, you will analyze the performance of a reference OpenMP implementation of this method, and then implement and analyze your own version using MPI.

The basic recurrence

At the heart of the method is the following basic recurrence. If l_{ij}^s represents the length of the shortest path from i to j that can be attained in at most 2^s steps, then

$$l_{ij}^{s+1} = \min_k \{l_{ik}^s + l_{kj}^s\}.$$

That is, the shortest path of at most 2^{s+1} hops that connects i to j consists of two segments of length at most 2^s , one from i to k and one from k to j . Compare this with the following formula to compute the entries of the square of a matrix A :

$$a_{ij}^2 = \sum_k a_{ik} a_{kj}.$$

These two formulas are identical, save for the niggling detail that the latter has addition and multiplication where the former has min and addition. But the basic pattern is the same, and all the tricks we learned when discussing matrix multiplication apply – or at least, they apply in principle. I’m actually going to be lazy in the implementation of `square`, which computes one step of this basic recurrence. I’m not trying to do any clever blocking. You may choose to be more clever in your assignment, but it is not required.

The return value for `square` is true if `l` and `lnew` are identical, and false otherwise.

```
int square(int n,                // Number of nodes
           int* restrict l,      // Partial distance at step s
           int* restrict lnew)   // Partial distance at step s+1
{
    int done = 1;
    #pragma omp parallel for shared(l, lnew) reduction(&& : done)
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < n; ++i) {
            int lij = lnew[j*n+i];
            for (int k = 0; k < n; ++k) {
```

```

        int lik = l[k*n+i];
        int lkj = l[j*n+k];
        if (lik + lkj < lij) {
            lij = lik+lkj;
            done = 0;
        }
    }
    lnew[j*n+i] = lij;
}
return done;
}

```

The value l_{ij}^0 is almost the same as the (i, j) entry of the adjacency matrix, except for one thing: by convention, the (i, j) entry of the adjacency matrix is zero when there is no edge between i and j ; but in this case, we want l_{ij}^0 to be “infinite”. It turns out that it is adequate to make l_{ij}^0 longer than the longest possible shortest path; if edges are unweighted, $n + 1$ is a fine proxy for “infinite.” The functions `infinitize` and `deinfinitize` convert back and forth between the zero-for-no-edge and $n + 1$ -for-no-edge conventions.

```

static inline void infinitalize(int n, int* l)
{
    for (int i = 0; i < n*n; ++i)
        if (l[i] == 0)
            l[i] = n+1;
}

static inline void deinfinitalize(int n, int* l)
{
    for (int i = 0; i < n*n; ++i)
        if (l[i] == n+1)
            l[i] = 0;
}

```

Of course, any loop-free path in a graph with n nodes can at most pass through every node in the graph. Therefore, once $2^s \geq n$, the quantity l_{ij}^s is actually the length of the shortest path of any number of hops. This means we can compute the shortest path lengths for all pairs of nodes in the graph by $\lceil \lg n \rceil$ repeated squaring operations.

The `shortest_path` routine attempts to save a little bit of work by only repeatedly squaring until two successive matrices are the same (as indicated by the return value of the `square` routine).

```

void shortest_paths(int n, int* restrict l)

```

```

{
    // Generate  $l_{ij}^{-0}$  from adjacency matrix representation
    infinitize(n, l);
    for (int i = 0; i < n*n; i += n+1)
        l[i] = 0;

    // Repeated squaring until nothing changes
    int* restrict lnew = (int*) calloc(n*n, sizeof(int));
    memcpy(lnew, l, n*n * sizeof(int));
    for (int done = 0; !done; ) {
        done = square(n, l, lnew);
        memcpy(l, lnew, n*n * sizeof(int));
    }
    free(lnew);
    deinfinitize(n, l);
}

```

The random graph model

Of course, we need to run the shortest path algorithm on something! For the sake of keeping things interesting, let's use a simple random graph model to generate the input data. The $G(n, p)$ model simply includes each possible edge with probability p , drops it otherwise – doesn't get much simpler than that. We use a thread-safe version of the Mersenne twister random number generator in lieu of coin flips.

```

int* gen_graph(int n, double p)
{
    int* l = calloc(n*n, sizeof(int));
    struct mt19937p state;
    sgenrand(10302011UL, &state);
    for (int j = 0; j < n; ++j) {
        for (int i = 0; i < n; ++i)
            l[j*n+i] = (genrand(&state) < p);
        l[j*n+j] = 0;
    }
    return l;
}

```

Result checks

Simple tests are always useful when tuning code, so I have included two of them. Since this computation doesn't involve floating point arithmetic, we

should get bitwise identical results from run to run, even if we do optimizations that change the associativity of our computations. The function `fletcher16` computes a simple [simple checksum](#) over the output of the `shortest_paths` routine, which we can then use to quickly tell whether something has gone wrong. The `write_matrix` routine actually writes out a text representation of the matrix, in case we want to load it into MATLAB to compare results.

```
int fletcher16(int* data, int count)
{
    int sum1 = 0;
    int sum2 = 0;
    for(int index = 0; index < count; ++index) {
        sum1 = (sum1 + data[index]) % 255;
        sum2 = (sum2 + sum1) % 255;
    }
    return (sum2 << 8) | sum1;
}

void write_matrix(const char* fname, int n, int* a)
{
    FILE* fp = fopen(fname, "w+");
    if (fp == NULL) {
        fprintf(stderr, "Could not open output file: %s\n", fname);
        exit(-1);
    }
    for (int i = 0; i < n; ++i) {
        for (int j = 0; j < n; ++j)
            fprintf(fp, "%d ", a[j*n+i]);
        fprintf(fp, "\n");
    }
    fclose(fp);
}
```

The main event

```
const char* usage =
    "path.x -- Parallel all-pairs shortest path on a random graph\n"
    "Flags:\n"
    "  - n -- number of nodes (200)\n"
    "  - p -- probability of including edges (0.05)\n"
    "  - i -- file name where adjacency matrix should be stored (none)\n"
    "  - o -- file name where output matrix should be stored (none)\n";

int main(int argc, char** argv)
```

```

{
    int n      = 200;           // Number of nodes
    double p = 0.05;           // Edge probability
    const char* ifname = NULL; // Adjacency matrix file name
    const char* ofname = NULL; // Distance matrix file name

    // Option processing
    extern char* optarg;
    const char* optstring = "hn:d:p:o:i:";
    int c;
    while ((c = getopt(argc, argv, optstring)) != -1) {
        switch (c) {
            case 'h':
                fprintf(stderr, "%s", usage);
                return -1;
            case 'n': n = atoi(optarg); break;
            case 'p': p = atof(optarg); break;
            case 'o': ofname = optarg; break;
            case 'i': ifname = optarg; break;
        }
    }

    // Graph generation + output
    int* l = gen_graph(n, p);
    if (ifname)
        write_matrix(ifname, n, l);

    // Time the shortest paths code
    double t0 = omp_get_wtime();
    shortest_paths(n, l);
    double t1 = omp_get_wtime();

    printf("== OpenMP with %d threads\n", omp_get_max_threads());
    printf("n:      %d\n", n);
    printf("p:      %g\n", p);
    printf("Time:    %g\n", t1-t0);
    printf("Check:  %X\n", Fletcher16(l, n*n));

    // Generate output file
    if (ofname)
        write_matrix(ofname, n, l);

    // Clean up
    free(l);
    return 0;
}

```